



**Institute for Software Integrated Systems
Vanderbilt University**



A Hands-on Introduction to Deep Learning with DeepForge

Brian Broll

brian.broll@vanderbilt.edu



Overview



- What is DeepForge?
 - Overview
 - Core Concepts
- CIFAR-10 Tutorial
- Advanced Topics
 - Infrastructure Integration
 - Data Provenance
 - Interactive Editing
- Demos
- Q&A

Questions welcome at any time!




Background



Deep Learning



- A *deep neural network* is an artificial neural network with multiple hidden layers
- Flexible enough to be applied to a number of problems including:
 - Speech-to-text
 - Image segmentation
 - Image classification
 - Learning embeddings
 - Styling images
- State of art in almost everything





DeepForge



What is DeepForge?



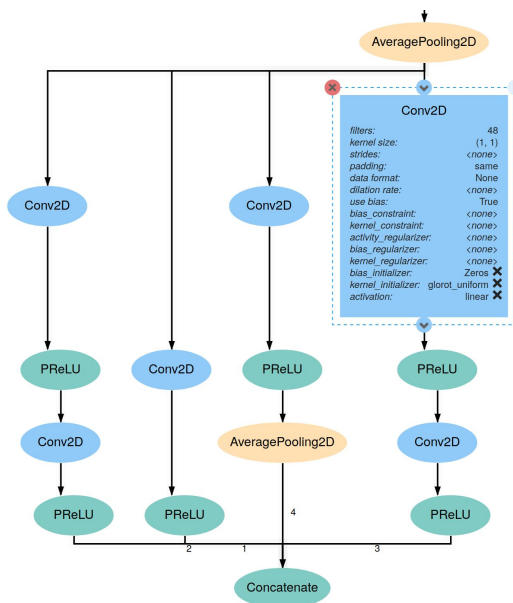
DeepForge is a platform for applying deep learning within diverse scientific domains that integrates with existing infrastructure.

The goal is to promote simplicity, collaboration, and reproducibility through:


- Domain specific visual/textual editors
 - Design neural networks (with error, shape feedback)
 - Create and execute workflows from the browser
- Automatic version control
- Data provenance
- Real-time collaborative editing
- Integration with existing HPC and storage resources



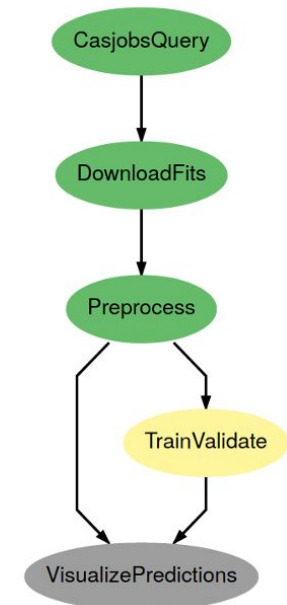
High-Level Overview



Design **Architectures**



Build Training/Testing **Pipelines**
(from **Operations**)



Execute Pipelines



Core Concepts



Core Concepts



- Two different types of concepts:
 - Concepts for creating executable pipelines
 - Concepts for designing neural network architectures
- Four main concepts are creating executable pipelines
 - Operations
 - Pipelines
 - Jobs
 - Executions
- Two high-level concepts for designing neural networks architectures
 - Architectures
 - Layers



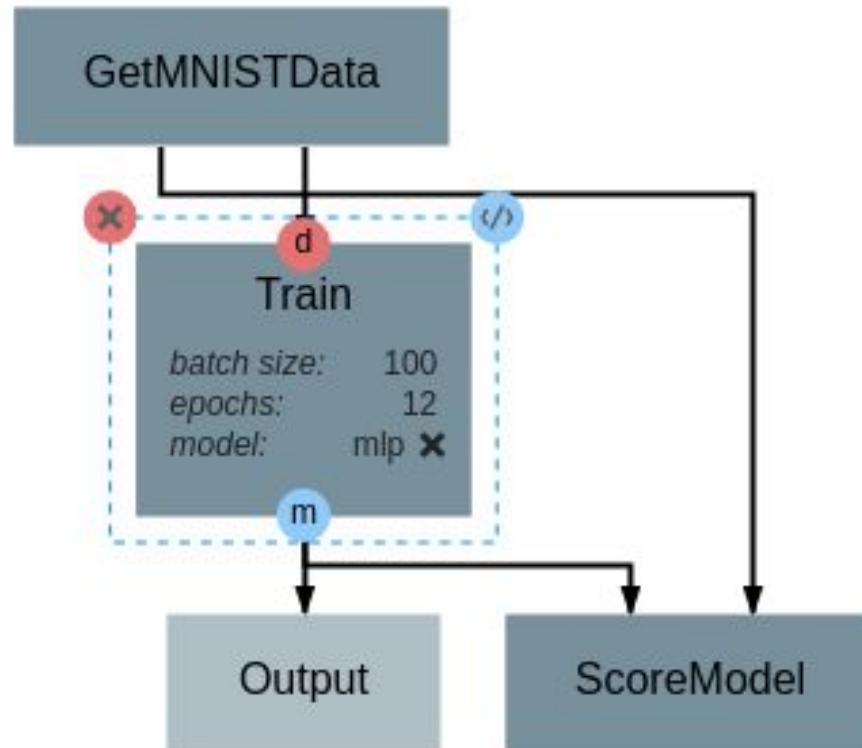
Pipeline Concepts



- **Operations** are functions with multiple, named inputs and outputs
- Attributes and references can be set at design time to specify configurable parameters
 - For example, *iterations* may be specified in a training operation
- **Pipelines** represent some machine learning task composed of operations
 - Examples include training, prediction, or data augmentation
- Pipelines can also contain *Input* and *Output* operations to specify inputs/outputs of the entire pipeline



Example Pipeline





Pipeline Concepts



- **Jobs** are executable operations which contain the operation definition and metadata about the execution
 - This includes console output and plots
- **Executions** represent an executable instance of a pipeline (composed of jobs)



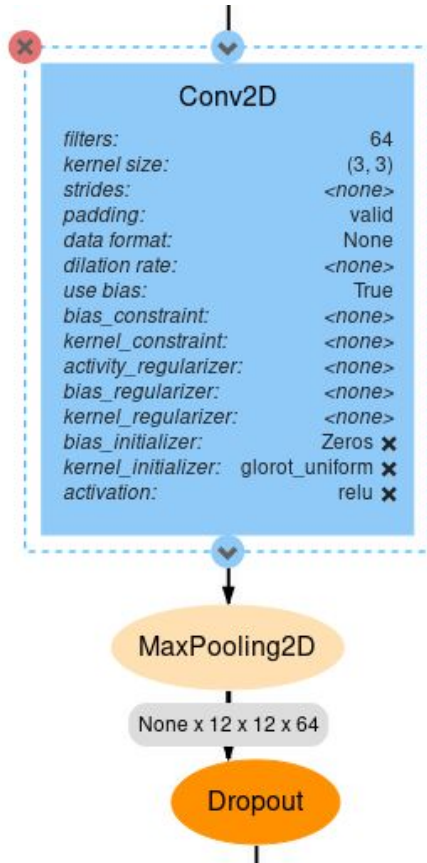
Neural Network Concepts



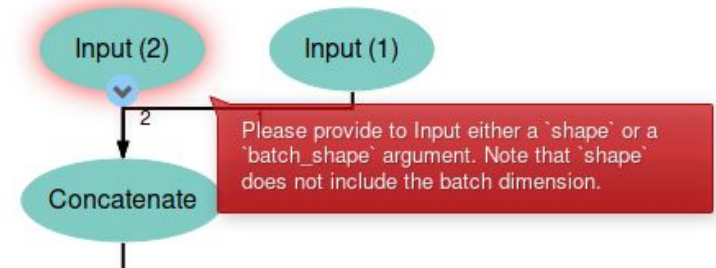
- The neural network concepts are pretty standard.
- Neural network architectures are composed of layers.
- Layers are parameterized in a variety of ways including the number of units, kernel size, dropout ratio, etc.



Neural Network Examples



Dimensionality Feedback



Error Messages
(During Construction)



Hands-on with CIFAR-10!



Advanced Topics



Infrastructure Integration



- DeepForge does not provide any compute or storage resources for use when executing pipelines*
- As shown in the tutorial, these capabilities are provided by integrating with existing platforms like SciServer.
- This integration is supported through the use of *compute adapters* and *storage adapters*.

* Technically, DeepForge can be configured to use a notebook-style of execution and storage but this is disabled on the public deployment.



Compute Adapters

- Compute adapters are used to execute jobs and pipelines. Each job will use the appropriate storage adapters for fetching and storing relevant data.
- Supported compute adapters:
 - SciServer Compute
 - WebGME Executor Framework (Connect your own machine)
 - Server execution (disabled on the public deployment)

Compute Options

Compute	<input type="text" value="SciServer Compute"/>
Username	<input type="text" value="deepforge"/> <small>SciServer username</small>
Password	<input type="password" value="*****"/> <small>SciServer password</small>
Compute Domain	<input type="text" value="Small Jobs Domain"/>

A small job shares resources with up to 4 other jobs and has a max quota for RAM of approx 32GB. A large job runs exclusively and has all CPU cores and RAM available (approx 240GB), however since only one large job will run at a time, there may be a longer wait for the job to start.



Storage Adapters



- Storage adapters are used to fetch and store data - DeepForge does not store any associated data (such as training data or trained models).
- Supported storage adapters:
 - SciServer Files
 - S3 Object Storage
 - WebGME Blob Storage (disabled on the public deployment)

Storage Options

Storage	<input type="text" value="SciServer Files Service"/>
Username	<input type="text" value="deepforge"/> <small>SciServer username</small>
Password	<input type="password" value="*****"/> <small>SciServer password</small>
Volume	<input type="text" value="deepforge/deepforge_data"/> <small>Volume to use for upload.</small>
Volume Pool	<input type="text" value="Storage"/> <small>Folders and files in User Volumes under "Storage" will be backed up and permanent, but there is a quota limit of 10GB. Folders and files in User Volumes under "Temporary" are not backed up, and will be deleted after a particular time period.</small>



Demos



Questions?



- Related Resources and Links:
 - Website: <https://deepforge.org>
 - Documentation: <http://deepforge.readthedocs.org/>
 - Deployment: <https://editor.deepforge.org>
 - Source Code:
<https://github.com/deepforge-dev/deepforge>
 - Slack Channel: <https://slack.deepforge.org>